




Figure 1: **INFERENCE and Alternative Approaches.** Vertical: GPU resources occupied or wasted by intercepted requests and other normal ones. Horizontal: timeline divided by iterations (dotted vertical lines). Hatch parts represent memory waste. INFERENCE adaptively combines preserve, our optimized discard, and our optimized swap. All green parts represent INFERENCE’s techniques.