

Cache Me if You Can: Repairing the KV State after Compression

Anonymous Authors¹

Abstract

As large language models (LLMs) support longer contexts and more capable agentic sessions, what matters in the accumulated context can change with each new query, tool output, test result, or browser observation. At the same time, inference still depends on a limited active key-value (KV) cache, and many existing KV-cache compression methods decide which past tokens stay active from the current prefix. As a result, tokens that matter in later turns may already have been evicted or compressed away in earlier compression decisions. We study post-compression KV cache repair, a runtime mechanism that restores previously evicted tokens to the active cache. REPAIRKV uses the interval between turns to reevaluate which tokens are important from evicted KV rows stored in a slower memory tier, then promotes a small subset before decoding resumes. **Controlled experiments show that REPAIRKV’s restoration of evicted tokens improves retrieval accuracy, persisting even across relevance changes and different initial eviction policies. A preliminary diagnostic on open-source repositories shows the same pattern: on Qwen2.5-7B-Instruct at 32K context, REPAIRKV achieves 91.0% retrieval on a four-query needle-in-a-haystack task versus 24.5% for the matched no-repair baseline at the same active-cache budget, with only 96 promoted tokens.**

1. Introduction

Modern language models support context lengths from tens of thousands to over a million tokens for retrieval-augmented and agentic workflows (??), but the KV cache holding past **token** activations grows linearly with sequence length, making it a primary GPU memory bottleneck **and**

an increasing runtime cost during decoding. KV-cache compression and offload (????) manage this cost by deciding which historical state remains active on GPU and which part is evicted or moved to a slower tier. This decision is usually made with information available at the current turn, and compression is often treated as a one-way operation: evict now, decode later from the retained active set.

This view misses a central feature of multi-turn inference: **new turn signals from** retrieved documents, tool outputs, test failures, or file changes can make previously low-priority context newly important. **These workflows also create** non-decoding intervals while the model waits for external results. Together, **the new signal and the pause create an opportunity to repair** the active KV state before the next decode, but **no existing work studies whether an already-compressed cache can be revised in runtime under a matched active-cache budget.**

This paper studies this gap. Our contributions are: (1) a matched resumed active-cache-budget protocol for testing whether repair helps without simply keeping more active tokens; (2) REPAIRKV, **an operator conditioned on the next turn’s signal that scores evicted KV rows and promotes a small budgeted subset back into the active cache before decoding resumes;** (3) controlled evidence across synthetic relevance shifts, different initial eviction policies, preliminary real-repository diagnostics, and runtime feasibility probes. On Qwen2.5-7B-Instruct at 32K context, 91.0% retrieval on a four-query needle-in-a-haystack task versus 24.5% for the matched no-repair baseline at the same active-cache budget, with only 96 promoted tokens.

2. Related Work

2.1. KV-cache compression and eviction

KV-cache compression and eviction methods reduce active KV state by retaining a fixed budget of tokens or blocks based on prefix-time, recency, or attention signals (?????), but this retention decision is usually one-way: rows dropped from the active set cannot re-enter when later turns make them important. This failure mode is observed in recent analyses of instruction and fact degradation under KV compression (?).

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

2.2. Query-aware retrieval and sparse attention

Query-aware methods such as QUEST emphasize the premise behind repair: token criticality depends on current information need, not only on the prefix that produced the compressed cache (?). Fine-grained KV retrieval systems use this signal during active decoding to select or fetch query-relevant rows from a retained or indexed KV store (????). Attention analyses likewise show that token importance can be query-dependent or recurrent rather than fixed once compressed (?). REPAIRKV builds on this query-dependence, but applies it at a different lifecycle point: after compression and before the next decode.

2.3. Long-context benchmarks

Long-context benchmarks stress whether models can use information across long inputs, from ordinary long documents to 100k-token settings, showing that application-visible context length does not remove the need for memory-management mechanisms (???). SCBench directly evaluates tasks around the KV-cache lifecycle itself (?). However, these benchmarks leave open what happens after a compression decision has already removed future-relevant tokens from the active cache.

2.4. KV serving and tiering

Serving systems address the systems side of long-context KV cache management. PagedAttention/vLLM improves GPU KV allocation and block management during serving (?), while prefix-caching systems reduce repeated prefill when requests share or resume prefixes (?). KV-tiering systems extend this idea across GPU, host memory, storage, and network tiers (???). Together, these systems determine where KV state lives and when it can be reused, but they do not address the complementary decision of which offloaded rows should re-enter the active cache for the next turn.

2.5. Multi-turn LLM serving

Multi-turn agent workloads pause decoding for external actions such as tool calls, web interaction, repository edits, and tests (????). Serving systems make these pauses operationally visible: Continuum schedules multi-turn agent requests around tool-call gaps and KV-cache lifetimes, while AgentServe targets local agent serving under tight GPU memory budgets (?). Pause-aware systems preserve paused-request KV to avoid repeated prefill (???). These same pauses often contain new relevance signals, creating the setting for repair.




Figure 1. REPAIRKV system overview. During agent decode pauses (tool calls, retrievals), a selected signal s — e.g., user input, tool output, retrieved content, recent generation — is used to rank host KV rows; the top- K are lifted back into GPU HBM at their original sequence positions.

3. REPAIRKV

3.1. Overview

Each compression step can only choose the active KV set with the information available at that point. REPAIRKV treats the resulting compressed state as revisable: when a new turn signal becomes available, the runtime can score offloaded evicted KV rows and promote a small subset before decoding resumes. Figure ?? summarizes this system view.

At the pause boundary, the abstraction has three objects:

- C_{base} : active evictable KV, with $|C_{\text{base}}| = B_{\text{base}}$ (base active-cache budget);
- W_N : offloaded evicted KV, hidden from decoding unless promoted;
- s : the pause boundary signal used to score evicted KV, such as a new user query, a tool result, the model’s recent generation, or other tokens that indicate upcoming significance.

The repair map chooses $S_K(s) \subseteq W_N$ with $|S_K| \leq K$, where K is the restore budget, and resumes from $C_{\text{repair}} = C_{\text{base}} \cup S_K$ at resumed active-cache budget $|C_{\text{repair}}| \leq B_{\text{base}} + K$.

3.2. Two-turn protocol

We isolate multi-turn repair with a protocol for the simplest controlled case: one relevance shift across two turns. A five-turn revisit diagnostic (Figure ??) confirms that the effect persists across repeated shifts.

At turn N , the model receives a long distractor document

followed by prompt Q_1 and generates the response to Q_1 from the full uncompressed cache. The runtime then compresses the distractor document KV to a base budget B_{base} on GPU and offloads the evicted rows to host memory as W_N . The Q_1 prompt, the generated response, and later Q_2 are preserved exactly. During the subsequent idle window, the turn- $(N+1)$ prompt Q_2 becomes known.

At turn $N+1$, the distractor document is not re-presented. REPAIRKV scores the evicted store W_N against the active post- Q_1 cache and the Q_2 signal, promotes a K -budgeted subset S_K back to GPU before Q_2 decoding, and the model answers Q_2 from C_{repair} .

The primary comparison excludes full-prefill recompute because it reprocesses the long context after Q_2 is known, while our comparison asks whether the already-compressed cache can be improved under the same resumed active-cache budget. The controlled prototype evaluates one paused request on one GPU with a single host-memory tier, a pre-materialized evicted-KV store, and the full Q_2 signal known before resumed decoding.

3.3. Repair operation

Intuitively, REPAIRKV **uses the next-turn signal to identify which evicted KV rows should be active before decoding resumes**. Algorithm ?? describes the general framework; the prototype is specified by two main choices: **SCORE** and **SELECT**, with $\mathcal{C} = C_{\text{base}}$ and $\mathcal{W} = W_N$.

For **SCORE**, post-RoPE Q_2 query tensors $\{q_{2,t}^{(\ell,h)}\}$ are **scored** against keys in the active and offloaded sets via

$$A^{(\ell,h)} = \text{softmax}\left(\frac{q_{2,:}^{(\ell,h)} [K_{C_{\text{base}}}, K_{W_N}]^\top}{\sqrt{d_k}}\right), \quad (1)$$

restricted to columns in W_N , then aggregated as $a_i = \text{mean}_{\ell,h} \max_t A_{t,i}^{(\ell,h)}$. Because the softmax normalizes over the active and offloaded keys, a_i is the attention mass row i would receive given C_{base} rather than a raw inner product. Ties break by the per-token score from the turn- N initial compression, then by ascending position.

For **SELECT**, each top-ranked anchor i is expanded into a local burst $\{i-L, \dots, i+R\} \cap (W_N \setminus S_K)$ under a strict $|S_K| \leq K$ budget. The right-biased $(L, R) = (2, 20)$ **reflects the local structure: high-scoring anchors often occur near queried keys, and answer-bearing value tokens usually follow nearby**. If anchor bursts do not exhaust K , remaining slots backfill from single ranked rows.

The current prototype scores token rows **fixed burst widths**. A systems implementation would promote cached blocks or pages.

Algorithm 1 General REPAIRKV Execution Framework

```

1: Input: turn signals  $\{s_N\}_{N=1}^T$ , eviction policy  $\mathcal{P}$ , base
   budget  $B_{\text{base}}$ , restore budget  $K$ .
2: Initialize: active cache  $\mathcal{C}$ , evicted store  $\mathcal{W} \leftarrow \emptyset$ .
3: for  $N = 1$  to  $T$  do
4:    $\text{ans}_N \leftarrow \text{DECODE}(\mathcal{C}, s_N)$  // Decoding
5:    $\mathcal{C}, \mathcal{W} \leftarrow \mathcal{P}(\mathcal{C}, \mathcal{W}, B_{\text{base}})$  // Compression
6:   obtain  $s_{N+1}$  // Idle window begins
7:   for  $i \in \mathcal{W}$  do
8:      $a_i \leftarrow \text{SCORE}(i \mid \mathcal{C}, s_{N+1})$  // Scoring
9:   end for
10:  rank  $\mathcal{W}$  by  $a_i$  // Ranking
11:   $S_K \leftarrow \text{SELECT}(\mathcal{W}, K), |S_K| \leq K$  // Selection
12:   $\mathcal{C} \leftarrow \mathcal{C} \cup S_K; \mathcal{W} \leftarrow \mathcal{W} \setminus S_K$  // Repair
13: end for
    
```

3.4. Matched-budget evaluation

All main claims compare REPAIRKV to a no-repair baseline at the same resumed active-cache budget. **In both conditions, Q_2 decoding sees the same number** of active evictable-context tokens in addition to the identical preserved Q_1 prompt and answer tail. The match is only on active GPU KV: extra offloaded rows, scorer time, and transfer time are not part of the matched budget; they are reported separately as service costs.

Let B_{match} denote the no-repair condition, which applies the base eviction policy with active evictable-context budget $B_{\text{base}} + K$. REPAIRKV reaches the same active budget by starting from B_{base} retained rows and promoting up to K rows from the host-memory store.

The main restore-budget frontier reports raw answer score. For effect size comparisons, we also report gain over matched no-repair,

$$\Delta_{\text{repair}} = \text{Score}(\text{RepairKV}) - \text{Score}(B_{\text{match}}) \quad (2)$$

where **Score** is the mean fraction of held-out Q_2 targets recovered correctly per example over the fixed evaluation set shared across all conditions. Stricter exact-set and precision-sensitive scoring serve as confirmatory checks.

4. Experimental Setup

4.1. Benchmark family

The main benchmark is a split-query version of multi-query needle-in-a-haystack (MQ-NIAH) derived from RULER (?). We use this family because it exposes the failure mode repair is meant to address: tokens that are unimportant for the first turn can become critical post-compression.

Each example contains key-value needles inside a long distractor context. The first turn queries one subset of keys,

the cache is compressed, and the second turn queries a different subset. This gives a controlled relevance shift with annotated future-relevant spans and a shared turn- N history across conditions.

We evaluate 2Q, 4Q, 6Q, and 8Q variants, which use the same task template but increase the number of queried key-value pairs. For each variant, we pool balanced partitions that avoid a recency shortcut: the keys requested in Q_2 are not the needles nearest to the end of the context. Appendix Figure ?? summarizes the exact splits.

4.2. Evaluation configuration

All main experiments use Qwen2.5-7B-Instruct (?) at 32K rendered context on a single RTX PRO 6000 Blackwell GPU. The initial post- Q_1 compression is context-only SnapKV compression (?): it is applied only to the evictable document context, while the Q_1 prompt and generated Q_1 response are preserved exactly.

Per-task base budgets are calibrated so matched no-repair is neither already near full-cache accuracy nor near zero, leaving room to measure whether repair improves the resumed cache. Unless stated otherwise, the fixed sweep is

- compressor: context-only SnapKV with $R_{\text{ctx}} = 128$ (SnapKV observation-window length);
- base budgets: $B_{\text{base}} = 8192$ for 2Q, $B_{\text{base}} = 16384$ for 4Q, and $B_{\text{base}} = 18432$ for 6Q/8Q;
- restore budgets: $K = 8, 16, 24, 32, 48, 64, 80, 96, 128$;
- replay: one shared full-cache-generated Q_1 response per setting and condition.

This replay removes turn- N generation differences, so the frontier isolates resumed-cache selection quality at turn $N+1$. MQ-NIAH-2Q/4Q/6Q use fixed 100-example evaluation sets per partition, while MQ-NIAH-8Q uses 24 examples per partition over five partitions. Prompt and decode settings are reported in the appendix.

4.3. Scorers

For the initial turn- N compression, we score context tokens using the generated Q_1 response rather than a generic last 32 token suffix. This makes the post- Q_1 compressed cache reflect the actual first turn computation.

For repair, the main quality curves use an exact Q_2 scorer. We append Q_2 to the compressed active cache, extract the model’s post-RoPE Q_2 query projections, and score active and offloaded key rows under those projections before selecting evicted rows for promotion. This isolates whether the newly revealed turn signal identifies rows worth restoring.

4.4. Baselines and references

We compare the full-cache reference A , the base compressed cache B before repair, matched no-repair B_{match} , Random- K and Oldest- K restore controls, and REPAIRKV. Random- K and Oldest- K promote from the same evicted KV store as REPAIRKV without scoring with Q_2 , testing whether gains come from query-aware repair or generic reinsertion.

The specificity study adds StaleQ- K (score with the previous turn query), WrongQ- K (score with another example’s Q_2) and Refresh-buffered, which reselects the whole $B_{\text{base}} + K$ context budget from active and offloaded rows after Q_2 is known, without full-prefix recomputation. Refresh-buffered serves to expose selector headroom rather than serving as the main baseline.

5. Results

5.1. Matched-budget frontier

Across MQ-NIAH-2Q/4Q/6Q/8Q, REPAIRKV separates from matched no-repair at moderate to large restore budgets. At $K = 96$, REPAIRKV scores 0.910 on 4Q, 0.989 on 6Q, and 0.960 on 8Q, while matched no-repair scores 0.245, 0.422, and 0.546, resp. Figure ?? reports the full raw Q_2 score frontier. Random- K and Oldest- K stay near matched no-repair, showing that generic reinsertion from the same evicted KV store does not explain the gains. MQ-NIAH-2Q saturates by $K = 80$, while matched no-repair remains near 0.06. Across all twelve balanced partitions at $K = 128$, REPAIRKV beats B_{match} and both content-agnostic controls. Even the weakest partition-level gains are 0.585, 0.387, and 0.292 for 4Q, 6Q, and 8Q, resp. (Appendix Figure ??).

5.2. Next-turn signal specificity

Repair gains depend on using the correct newly revealed signal. At $K = 48$, StaleQ- K and WrongQ- K stay near matched no-repair, with gains 0.028 $[-0.021, 0.083]$ and 0.028 $[-0.014, 0.069]$, while REPAIRKV reaches 0.326 $[0.243, 0.410]$. Thus, the improvement is not explained by scoring with any query-like text; the signal must match the upcoming turn. Refresh-buffered reaches 1.000, showing that relevant rows remaining available for stronger selectors to exploit, motivating an open research direction.

5.3. Repeated relevance shifts

The repair effect persists when relevance changes repeatedly. On a fixed five-turn MQ-NIAH-8Q revisit schedule ($K = 80$, $n = 24$), REPAIRKV gains 0.542 over matched no-repair on non-initial turns (95% CI $[0.458, 0.620]$). Random- K and Oldest- K gain only 0.010 and 0.021, and REPAIRKV remains above StaleQ- K , supporting that repair




Figure 2. Raw exact Q_2 score under matched resumed-cache budgets. Solid curves are REPAIRKV. Dotted curves are matched no-repair. Right labels report query count and Δ_{96} , the score gain over matched no-repair at $K = 96$.

Table 1. Next-turn signal specificity at $K = 48$ on MQ-NIAH-4Q (72 paired example-partition cases). Stale and donor controls test whether any query-like signal suffices; Refresh is a full-budget reselector over active and offloaded KV.

Condition	Q_2 score	Δ vs. matched [95% CI]
Matched no-repair	0.215	—
StaleQ- K	0.243	+0.028 [−0.021, 0.083]
WrongQ- K	0.243	+0.028 [−0.014, 0.069]
REPAIRKV	0.542	+0.326 [0.243, 0.410]
Refresh-buffered	1.000	+0.785 [0.722, 0.847]

can track the current turn signal across repeated controlled shifts.

5.4. Eviction-policy sensitivity

REPAIRKV’s improvement improves caches produced by other first-stage eviction policies as well. On MQ-NIAH-4Q, we replace the primary SnapKV-style compressor with H₂O-style accumulated attention retention and a StreamingLLM-style sink-plus-recent rule. In both cases, REPAIRKV improves over each policy’s matched no-repair baseline while Random- K and Oldest- K remain near baseline (Figure ??). Under the StreamingLLM-style policy, for example, REPAIRKV gains 0.431 at $K = 128$. A separate Scissorhands-style persistence stress test also passes the locked gate (Appendix Table ??).

5.5. Real-repository diagnostic

To test external validity beyond MQ-NIAH, we ran a controlled relevance shift diagnostic on 48 callsite examples from 12 open-source repositories drawn from the SWE-




Figure 3. Five-turn relevance-shift diagnostic on MQ-NIAH-8Q ($B_{\text{base}}=18,432$, $K=80$, $n=24$). Groups show post-priming exact scores at equal active-cache budget. Intervals are bootstrap 95% CIs for REPAIRKV, and Stale-Q uses the previous-turn query.




Figure 4. First-stage eviction-policy sensitivity on MQ-NIAH-4Q ($n=24$, $B_{\text{base}}=16,384$). Curves show exact Q_2 gain over each policy’s matched no-repair baseline. The gray band spans Random- K /Oldest- K gains.

bench pool (?) at fixed commits. Each example loads 32K tokens of line-numbered repository context. The first turn asks about a file unrelated to the answer, encouraging compression to evict the later relevant code. The second turn asks for a target identifier at a callsite.

Between turns, an *event* arrives: a description of the callsite with the target identifier redacted. This event replaces the synthetic Q_2 signal used in MQ-NIAH, and all conditions decode the same event and question prompt. At $K = 192$, event only REPAIRKV improves exact identifier accuracy from 18.8% matched no-repair to 72.9%. Non-query controls and ToolFile- K remain near matched (Table ??). Thus, the gain comes from using the event as a relevance signal, not from generic reinsertion or simply knowing the file name.

Label-assisted references show remaining headroom: File gated REPAIRKV reaches 83.3%, and AnchorWindow- K reaches 89.6%. This is a small diagnostic rather than a full benchmark, but it shows some initial evidence that the repair approach improves quality beyond the synthetic MQ-NIAH setting. Appendix Figure ?? gives the full audit.

Table 2. Preliminary real-repository diagnostic at $K = 192$ over 48 callsite cases from 12 open-source repositories (SWE-bench pool, fixed commits). Accuracy is exact target-identifier recovery. Non-query summarizes the best of Random- K , Oldest- K , StaleQ- K , and WrongQ- K analogues. AnchorWindow and File-gated are label-assisted; non-label methods are Matched, RepairKV, and ToolFile- K .

Condition	Accuracy
Matched no-repair	18.8%
Best non-query control	16.7%
ToolFile- K	18.8%
REPAIRKV	72.9%
File-gated REPAIRKV	83.3%
AnchorWindow- K	89.6%

5.6. Runtime

We run a separate GPU runtime probe for the repair operation, measuring chunked scan, top- K selection, and KV movement over pinned host-memory BF16 tensors with Qwen-shaped dimensions. At $K = 5000$, p95 repair latency is 50.0 ms at 32K offloaded candidates, 1.20 s at 1M, and 4.64 s at 4M (Figure ??). These numbers are hardware- and model-dependent: larger KV dimensions increase the constant factor, while newer deployment GPUs may reduce it. The key property is linearly scaling in the number of candidate KV rows, which keeps repair in the seconds range even at million-row candidate pools, unlike approaches that recompute or score over full context-query interactions with quadratic cost.

These costs are small relative to many agent and tool-use pauses: OS-level execution, including tool calls and setup, can account for 56–74% of task latency (?), and tool calls range from sub-second file/shell operations to multi-second tests or code execution (?). More broadly, the low repair cost suggests a runtime role beyond using otherwise idle GPU time: repair can be run opportunistically at turn boundaries, tool results, retrieval updates, or other state changes to trade a small amount of memory movement and time for substantially better resumed-cache quality.

6. Conclusion

REPAIRKV makes the case for revisable KV compression. Existing compression policies decide which past tokens remain active using the information available at compression time. Our results show that this decision can be usefully revised when a later turn reveals a new relevance signal. Across MQ-NIAH-2Q/4Q/6Q/8Q, repeated relevance shifts, alternative eviction policies, and a small real-repository diagnostic, promoting a small set of relevant evicted rows substantially improves resumed-cache quality under the same active-cache budget.




Figure 5. Runtime-capacity probe on the evaluation GPU. (a) Component-summed p95 service time for chunked scan, top- K selection, and KV movement at query length 64 over pinned host-memory BF16 KV-shaped tensors. (b) Stage-wise p95 component time for $K = 5000$. Generation and answer evaluation are excluded.

The broader direction is mutable KV memory for resource adaptive long-context inference. Repair suggests that runtimes need not treat compression as a one way loss: they can maintain searchable evicted state, score it when new signals arrive, and use cheap promotion to trade a small amount of memory movement for better response quality. This opens several follow-up directions: stronger Q_2 -aware selectors, page-level repair policies, scheduler-aware repair at arbitrary turn or tool events, and benchmarks with production-style multi-turn traces, hidden future relevance, and explicit accounting for active-cache, offloaded storage, transfer, and recompute cost.

Impact Statement

This paper studies inference-time KV-cache repair for long-context language models. Its intended impact is to reduce GPU memory pressure while improving response quality in long-context and agentic settings where relevance changes between turns. We do not foresee novel ethical risks beyond those already inherent to large language model use.

References

- Abhyankar, R., He, Z., Srivatsa, V., Zhang, H., and Zhang, Y. InferCept: Efficient intercept support for augmented large language model inference. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.

- Bai, Y., Lv, X., Zhang, J., Lyu, H., Tang, J., Huang, Z., Du, Z., Liu, X., Zeng, A., Hou, L., Dong, Y., Tang, J., and Li, J. LongBench: A bilingual, multitask benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- Bian, S., Yan, M., Jayarajan, A., Pekhimenko, G., and Venkataraman, S. What limits agentic systems efficiency. In *NeurIPS 2025 Workshop on Scaling Environments for Agents (SEA)*, 2025.
- Bian, Z., Wu, F., Ma, T., and Zhuo, Y. Tokencake: A KV-cache-centric serving framework for LLM-based multi-agent applications. *arXiv:2510.18586*, 2025.
- Chen, A., Geh, R., Grover, A., Van den Broeck, G., and Israel, D. The pitfalls of KV cache compression. *arXiv:2510.00231*, 2025.
- Dzikanyanga, G., Yang, W., Huang, H., Wu, D., Wang, S., Xia, W., and K C, S. TTKV: Temporal-tiered KV cache for long-context LLM inference. *arXiv:2604.19769*, 2026.
- Gao, B., He, Z., Sharma, P., Kang, Q., Jevdjic, D., Zhao, J., and Zuo, P. Cost-efficient large language model serving for multi-turn conversations with CachedAttention. In *Proceedings of the 2024 USENIX Annual Technical Conference (USENIX ATC)*, 2024.
- Hsieh, C.-Y., Sun, S., Krizan, S., Acharya, S., Rekesh, D., Jia, F., Zhang, Y., and Ginsburg, B. RULER: What’s the real context window size of your LLM. In *Conference on Language Modeling (COLM)*, 2024.
- Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. SWE-bench: Can language models resolve real-world GitHub issues. In *International Conference on Learning Representations (ICLR)*, 2024.
- Kwon, W., Li, Z., Zhuang, S., Sheng, Y., Zheng, L., Yu, C. H., Gonzalez, J. E., Zhang, H., and Stoica, I. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th ACM Symposium on Operating Systems Principles (SOSP)*, 2023.
- Li, Y., Huang, Y., Yang, B., Venkitesh, B., Locatelli, A., Ye, H., Cai, T., Lewis, P., and Chen, D. SnapKV: LLM knows what you are looking for before generation. *arXiv:2404.14469*, 2024.
- Li, Z., Yang, Y., Zhang, Y., and Liu, Z. PyramidKV: Dynamic KV cache compression based on pyramidal information funneling. *arXiv:2406.02069*, 2024.
- Li, Y., Jiang, H., Wu, Q., Luo, X., Ahn, S., Zhang, C., Abdi, A. H., Li, D., Gao, J., Yang, Y., and Qiu, L. SCBench: A KV cache-centric analysis of long-context methods. In *International Conference on Learning Representations (ICLR)*, 2025.
- Li, H., Mang, Q., He, R., Zhang, Q., Mao, H., Chen, X., Zhou, H., Cheung, A., Gonzalez, J., and Stoica, I. Continuum: Efficient and robust multi-turn LLM agent scheduling with KV cache time-to-live. *arXiv:2511.02230*, 2026.
- Liu, Y., Cheng, Y., Yao, J., An, Y., Chen, X., Feng, S., Huang, Y., Shen, S., Zhang, R., Du, K., and Jiang, J. LMCache: An efficient KV cache layer for enterprise-scale LLM inference. *arXiv:2510.09665*, 2025.
- Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., Gu, Y., Ding, H., Men, K., Yang, K., Zhang, S., Deng, X., and others. AgentBench: Evaluating LLMs as agents. In *International Conference on Learning Representations (ICLR)*, 2024.
- Liu, Z., Desai, A., Liao, F., Wang, W., Xie, V., Xu, Z., Kyrillidis, A., and Shrivastava, A. Scissorhands: Exploiting the persistence of importance hypothesis for LLM KV cache compression at test time. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Qin, Y., Liang, S., Ye, Y., Zhu, K., Yan, L., Lu, Y., Lin, Y., Cong, X., Tang, X., Qian, B., Zhao, S., Hong, L., and others. ToolLLM: Facilitating large language models to master 16000+ real-world APIs. In *International Conference on Learning Representations (ICLR)*, 2024.
- Qwen Team. Qwen2.5 technical report. *arXiv:2412.15115*, 2024.
- Pagonas, N., Chung, Y., Kaffes, K., and Krishnamurthy, A. Cortex: Workflow-aware resource pooling and scheduling for agentic serving. *arXiv:2510.14126*, 2025.
- Qi, Y., Chen, X., Jiang, H., Wang, Q., Peng, B., and Papanas, T. ParisKV: Fast and drift-robust KV-cache retrieval for long-context LLMs. *arXiv:2602.07721*, 2026.
- Wu, W., Wang, Y., Xiao, G., Peng, H., and Fu, Y. Retrieval heads mechanistically explain long-context factuality. *arXiv:2404.15574*, 2024.
- Sun, H., Chang, L.-W., Bao, W., Zheng, S., Zheng, N., Liu, X., Dong, H., Chi, Y., and Chen, B. ShadowKV: KV cache in shadows for high-throughput long-context LLM inference. *arXiv:2410.21465*, 2025.
- Tang, J., Zhao, Y., Zhu, K., Xiao, G., Kasikci, B., and Han, S. QUEST: Query-aware sparsity for efficient long-context LLM inference. In *Proceedings of the 41st International Conference on Machine Learning (ICML)*, 2024.

- Wang, D., Liu, Z., Wang, S., Ren, Y., Deng, J., Hu, J., Chen, T., and Yang, H. FIER: Fine-grained and efficient KV cache retrieval for long-context LLM inference. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.
- Xiao, G., Tang, Y., Zuo, J., Guo, J., Yang, S., Tang, H., Fu, Y., and Han, S. Efficient streaming language models with attention sinks. In *International Conference on Learning Representations (ICLR)*, 2024.
- Zhang, Z., Sheng, Y., Zhou, T., Chen, T., Zheng, L., Cai, R., Song, Z., Tian, Y., Ré, C., Barrett, C., Wang, Z., and Chen, B. H₂O: Heavy-hitter oracle for efficient generative inference of large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Zhang, X., Chen, Y., Hu, S., Xu, Z., Chen, J., Hao, M. K., Han, X., Thai, Z. L., Wang, S., Liu, Z., and Sun, M. ∞ Bench: Extending long context evaluation beyond 100K tokens. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.
- Zhang, H., Zhang, H., Ma, X., Zhang, J., and Guo, S. LazyEviction: Lagged KV eviction with attention pattern observation for efficient long reasoning. *arXiv:2506.15969*, 2025.
- Zhang, Y., Yan, Y., Yang, N., and Yuan, D. AgentServe: Algorithm-system co-design for efficient agentic AI serving on a consumer-grade GPU. *arXiv:2603.10342*, 2026.
- Zheng, Y., Fan, J., Fu, Q., Yang, Y., Zhang, W., and Quinn, A. AgentCgroup: Understanding and controlling OS resources of AI agents. *arXiv:2602.09345*, 2026.
- Zhou, S., Xu, F. F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Ou, T., Bisk, Y., Fried, D., Alon, U., and Neubig, G. WebArena: A realistic web environment for building autonomous agents. In *International Conference on Learning Representations (ICLR)*, 2024.

A. Additional Evaluation Details

Protocol details. Concise split prompts request values only, comma-separated, with a short answer prefix and fixed decode budgets of 24 tokens for MQ-NIAH-4Q, 48 tokens for MQ-NIAH-6Q, and 64 tokens for MQ-NIAH-8Q. Generation is deterministic greedy decoding with no sampling. The main examples use dataset seed offset 0. Random- K uses a deterministic example- and K -dependent seed. Bootstrap intervals are percentile intervals over paired example-partition rows with fixed seeds. Main Qwen runs use BF16 weights on one RTX PRO 6000 Blackwell GPU with Python 3.12.3, PyTorch 2.10.0+cu128, and Transformers 5.2.0. The exact- Q_2 attention-query selector appends Q_2 to the compressed active cache, extracts the model’s post-RoPE query projections for the Q_2 prompt tokens, and scores active and offloaded key rows under those projections before selecting evicted rows for promotion. This selector is distinct from decode-time answer-token attention and from the final answer metric. The Q_2 -appended-state proxy appends Q_2 to the compressed cache and uses the appended cache rows as a second Q_2 -conditioned scoring signal.

Span-group diagnostic. SpanRef- K starts from the same B state and offloaded evicted-KV store as REPAIRKV. It enumerates the small family of annotated Q_2 answer-span-group subsets, one key/value span group per requested needle, evaluates each subset once under actual Q_2 generation, and chooses the highest-scoring subset with total cost at most K . It is exact over this benchmark-specific span-group family, but not over all arbitrary token subsets or possible final generations. Because the selected span-group subset may use fewer than K rows, a repair selector can occasionally score higher at low budgets by spending the full budget on useful local neighborhoods around future-relevant spans. Thus SpanRef- K is a diagnostic span-group reference, not a ceiling.

Partition choices. For MQ-NIAH-4Q, the pooled panel uses 14→23, 24→13, and 34→12. These are the balanced 2|2 partitions whose Q_2 side excludes the tail-anchored fourth needle. For MQ-NIAH-6Q, the pooled panel uses 156→234, 256→134, 356→124, and 456→123, so the second turn excludes the two latest needles. For MQ-NIAH-8Q, the pooled panel uses 5678→1234, 1678→2345, 2678→1345, 3678→1245, and 4678→1235, again excluding the two latest needles from the second turn. Within-turn order is not treated as a separate condition because the prompt asks for comma-separated values only and scoring checks whether each requested value is recovered. Recency-favorable splits such as 12→34 are left to separate diagnostics outside the main future-query repair result.

Scissorhands-style persistence stress test. Table ?? repeats the MQ-NIAH-6Q locked protocol after replacing the initial SnapKV-style eviction policy with a one-shot persistence policy inspired by Scissorhands. The test asks whether repair still works when the compressed active cache is chosen by past attention persistence in place of the primary eviction policy. Canonical online per-head Scissorhands reproduction remains future work.

Table 3. Scissorhands-style persistence stress test on MQ-NIAH-6Q ($B_{\text{base}}=18,432$, $n=24$). Values are exact scores. Best ctrl. is the stronger of Random- K and Oldest- K ; the interval is the paired bootstrap 95% CI for REPAIRKV- B_{match} .

K	Full	B_{match}	Best ctrl.	REPAIRKV	Gain CI
48	0.990	0.365	0.382	0.712	[0.292, 0.403]
64	0.990	0.365	0.382	0.795	[0.378, 0.490]
80	0.990	0.378	0.385	0.934	[0.503, 0.604]
96	0.990	0.375	0.382	0.993	[0.566, 0.670]
128	0.990	0.368	0.378	0.990	[0.573, 0.674]

B. Additional Discussion

Tiered-cache scaling. Repair creates a two-level cache-budget problem because the main curves match the resumed active GPU KV budget while tracking total host-plus-device storage separately. A strict device-resident policy chooses the active-cache budget, while a looser off-device policy decides which evicted rows remain searchable, compressed, summarized, or recomputable. In the Qwen-shaped accounting above, 4M searchable rows already imply about 240 GB before metadata. Thus REPAIRKV supports the middle operation in a tiered hierarchy: turn-conditioned promotion from host memory during an idle window. This is a hardware-facing implication; hardware measurement remains future work.

Idle-window repair interface. A practical implementation would expose a host-memory index over evicted KV units, a time-budgeted selector conditioned on the next-turn signal, asynchronous promotion into the active cache, and fallbacks to

no-repair or full recompute. This turns a pause into a schedulable resource: the system can spend only the slack that remains after queueing, transfer, and generation constraints, and can move the selector from token rows to pages or compressed metadata as hardware and serving constraints require.

Preliminary nonlinear variable-tracking diagnostic. As an exploratory stress test, we also tested a modified RULER variable-tracking task that makes the dependency structure less linear than the standard 4-hop chain. The diagnostic uses an 8-hop permuted chain with two distractor divergences: some assignments introduce branches that are not used by the final answer, so the model must recover the relevant path through a small information graph rather than follow a tail-friendly linear sequence. In 32K-context Qwen2.5-7B-Instruct runs, SnapKV compression degraded this task sharply even at moderate retained-cache budgets. Preliminary repair runs recovered substantial accuracy from the same compressed states using small restore fractions. These results are not part of the main evidence, but they motivate a broader benchmark for nonlinear and conflicting relevance patterns, where eviction policies may fail even when standard long-context tasks appear solved. Turning this observation into a robust benchmark is a natural next step for testing when compression and repair need to operate together.

Future benchmark axes. Future benchmarks should vary the number, timing, and conflict structure of new-turn constraints. The 2Q and 8Q breadth checks suggest that easy regimes can saturate while stricter multi-query turns expose how quickly eviction failures accumulate. The five-turn diagnostic strengthens the dynamic-workflow evidence, but it remains synthetic. Broader task families should keep the same matched-budget controls while adding real tool traces and stronger selector families.

C. Supplementary Experimental Views

The supplementary figures probe external validity, query-count and operating-regime breadth, partition consistency, runtime capacity, and appendix selector diagnostics. Exact audit values remain in released result artifacts.




Figure 6. Controlled real-repository relevance-shift diagnostic over 48 callsite examples from 12 repositories drawn from the SWE-bench repository pool at pinned commits. Repair sees only the event cue. All conditions decode the same event-plus-question prompt. Left: exact identifier accuracy at $K = 96$ and $K = 192$. Right: $K = 192$ REPAIRKV gain over matched no-repair after removing cue-only or answer-retention rows. ToolFile- K is file-name-assisted with oldest-row backfill. File-gated REPAIRKV restricts candidates to the event-named file before applying the event-only scorer. Lexical is an answer-sanitized diagnostic outside the deployable-selector comparison. AnchorWindow- K is a label-assisted locality reference outside the deployable-baseline and end-to-end issue-resolution comparisons.




Figure 7. Supplementary breadth diagnostics. (a) Cells show REPAIRKV score gain over matched no-repair for $K = 48$ and $K = 96$ across 2-, 3-, 4-, 6-, and 8-question split-query variants. (b) Cells show the fraction of available SpanRef- K diagnostic gain recovered by REPAIRKV, clipped to $[0, 1]$, for MQ-NIAH-4Q/6Q. Saturated cells with little matched-to-SpanRef- K gap should not be read as upper-bound statements.




Figure 8. Per-partition endpoint scores at $K = 128$ on the final exact evaluations. Gray segments show the change from matched no-repair to REPAIRKV, orange points show the benchmark-metadata SpanRef- K diagnostic, and labels report the REPAIRKV difference over matched no-repair.




Figure 9. Mechanism diagnostic for the final exact MQ-NIAH evaluations. Curves show overlap between promoted tokens and annotated Q_2 answer spans. Content-agnostic promotions stay near baseline while REPAIRKV increasingly covers future-relevant spans.




Figure 10. Supplementary runtime and selector diagnostics. (a) Synthetic move-and-inject capacity for Qwen-shaped BF16 KV on the evaluation GPU, separate from the Figure ?? scan/select envelope. Measurements use pinned host-memory promoted blocks and report p95 latency as K or active cache size grows. (b) Appendix proxy-selector diagnostic on MQ-NIAH-4Q/6Q (100 examples per split). At $K = 96$, content-agnostic promotions stay within 0.010/0.004 of matched no-repair while proxy REPAIRKV remains well above matched.